

Starbucks_Capstone_notebook

March 30, 2022

1 Starbucks Capstone Challenge

1.0.1 Introduction

This data set contains simulated data that mimics customer behavior on the Starbucks rewards mobile app. Once every few days, Starbucks sends out an offer to users of the mobile app. An offer can be merely an advertisement for a drink or an actual offer such as a discount or BOGO (buy one get one free). Some users might not receive any offer during certain weeks.

Not all users receive the same offer, and that is the challenge to solve with this data set.

Your task is to combine transaction, demographic and offer data to determine which demographic groups respond best to which offer type. This data set is a simplified version of the real Starbucks app because the underlying simulator only has one product whereas Starbucks actually sells dozens of products.

Every offer has a validity period before the offer expires. As an example, a BOGO offer might be valid for only 5 days. You'll see in the data set that informational offers have a validity period even though these ads are merely providing information about a product; for example, if an informational offer has 7 days of validity, you can assume the customer is feeling the influence of the offer for 7 days after receiving the advertisement.

You'll be given transactional data showing user purchases made on the app including the timestamp of purchase and the amount of money spent on a purchase. This transactional data also has a record for each offer that a user receives as well as a record for when a user actually views the offer. There are also records for when a user completes an offer.

Keep in mind as well that someone using the app might make a purchase through the app without having received an offer or seen an offer.

1.0.2 Example

To give an example, a user could receive a discount offer buy 10 dollars get 2 off on Monday. The offer is valid for 10 days from receipt. If the customer accumulates at least 10 dollars in purchases during the validity period, the customer completes the offer.

However, there are a few things to watch out for in this data set. Customers do not opt into the offers that they receive; in other words, a user can receive an offer, never actually view the offer, and still complete the offer. For example, a user might receive the "buy 10 dollars get 2 dollars off offer", but the user never opens the offer during the 10 day validity period. The customer spends 15 dollars during those ten days. There will be an offer completion record in the data set; however, the customer was not influenced by the offer because the customer never viewed the offer.

1.0.3 Cleaning

This makes data cleaning especially important and tricky.

You'll also want to take into account that some demographic groups will make purchases even if they don't receive an offer. From a business perspective, if a customer is going to make a 10 dollar purchase without an offer anyway, you wouldn't want to send a buy 10 dollars get 2 dollars off offer. You'll want to try to assess what a certain demographic group will buy when not receiving any offers.

1.0.4 Final Advice

Because this is a capstone project, you are free to analyze the data any way you see fit. For example, you could build a machine learning model that predicts how much someone will spend based on demographics and offer type. Or you could build a model that predicts whether or not someone will respond to an offer. Or, you don't need to build a machine learning model at all. You could develop a set of heuristics that determine what offer you should send to each customer (i.e., 75 percent of women customers who were 35 years old responded to offer A vs 40 percent from the same demographic to offer B, so send offer A).

2 Data Sets

The data is contained in three files:

- `portfolio.json` - containing offer ids and meta data about each offer (duration, type, etc.)
- `profile.json` - demographic data for each customer
- `transcript.json` - records for transactions, offers received, offers viewed, and offers completed

Here is the schema and explanation of each variable in the files:

portfolio.json * `id` (string) - offer id * `offer_type` (string) - type of offer ie BOGO, discount, informational * `difficulty` (int) - minimum required spend to complete an offer * `reward` (int) - reward given for completing an offer * `duration` (int) - time for offer to be open, in days * `channels` (list of strings)

profile.json * `age` (int) - age of the customer * `became_member_on` (int) - date when customer created an app account * `gender` (str) - gender of the customer (note some entries contain 'O' for other rather than M or F) * `id` (str) - customer id * `income` (float) - customer's income

transcript.json * `event` (str) - record description (ie transaction, offer received, offer viewed, etc.) * `person` (str) - customer id * `time` (int) - time in hours since start of test. The data begins at time `t=0` * `value` - (dict of strings) - either an offer id or transaction amount depending on the record

Note: If you are using the workspace, you will need to go to the terminal and run the command `conda update pandas` before reading in the files. This is because the version of pandas in the workspace cannot read in the `transcript.json` file correctly, but the newest version of pandas can. You can access the terminal from the orange icon in the top left of this notebook.

You can see how to access the terminal and how the install works using the two images below. First you need to access the terminal:

Then you will want to run the above command:

Finally, when you enter back into the notebook (use the jupyter icon again), you should be able to run the below cell without any errors.

```

In [1]: import pandas as pd
import numpy as np
import math
import json
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
from joblib import dump, load

from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, GradientDescentClassifier
from sklearn.neighbors import KNeighborsClassifier as KNN
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report, f1_score
from sklearn import preprocessing

from sklearn.model_selection import train_test_split, RandomizedSearchCV, GridSearchCV
from sklearn.metrics import accuracy_score, f1_score, fbeta_score, make_scorer, mean_squared_error
from sklearn.linear_model import LogisticRegression

In [2]: # read in the json files
portfolio = pd.read_json('data/portfolio.json', orient='records', lines=True)
profile = pd.read_json('data/profile.json', orient='records', lines=True)
transcript = pd.read_json('data/transcript.json', orient='records', lines=True)

```

3 Data Cleaning

In order to have a good understanding of the data, let us take time to clean and explore the data to determine the best procedure to clean up the data. This would enable us check for missing value, visualizing the data distribution

```
In [3]: portfolio.head()
```

```

Out[3]:
   channels  difficulty  duration \
0  [email, mobile, social]      10      7
1  [web, email, mobile, social]    10      5
2      [web, email, mobile]        0      4
3      [web, email, mobile]        5      7
4      [web, email]              20     10

   id      offer_type  reward
0  ae264e3637204a6fb9bb56bc8210ddfd      bogo      10
1  4d5c57ea9a6940dd891ad53e9dbe8da0      bogo      10
2  3f207df678b143eea3cee63160fa8bed  informational      0
3  9b98b8c7a33c4b65b9aebfe6a799e6d9      bogo      5
4  0b1e1539f2cc45b7b9fa7c272da2e1d7      discount      5

```

```
In [27]: profile.head()
```

```
Out[27]:
```

	age	became_member_on	gender	id	income	\
0	118	1872	None	68be06ca386d4c31939f3a4f0e3dd783	NaN	
1	55	1719	F	0610b486422d4921ae7d2bf64640c50b	112000.0	
2	118	1357	None	38fe809add3b4fcf9315a9694bb96ff5	NaN	
3	75	1786	F	78afa995795e4d85b5d9ceeca43f5fef	100000.0	
4	118	1699	None	a03223e636434f42ac4c3df47e8bac43	NaN	

	valid
0	0
1	1
2	0
3	1
4	0

```
In [5]: transcript.head()
```

```
Out[5]:
```

	event	person	time	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0	
1	offer received	a03223e636434f42ac4c3df47e8bac43	0	
2	offer received	e2127556f4f64592b11af22de27a7932	0	
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0	
4	offer received	68617ca6246f4fbc85e91a2a49552598	0	

	value
0	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}
1	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}
2	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}
3	{'offer id': 'fafdc668e3743c1bb461111dcafc2a4'}
4	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}

```
In [6]: from collections import Counter
```

```
def clean_portfolio(df):
```

```
    """
```

```
    Description:
```

```
        Deserialize the channels attribute to create column header for each of them along with their respective values
```

```
        Merge the portfolio dataframe with the stacked data using the id field
```

```
        Drop the channels column from the dataset
```

```
        Rename the id field as offer_id
```

```
        Convert the duration data to hours instead of days
```

```
        drop any duplicates if any
```

```
    Input:
```

```
        - df: original portfolio dataframe
```

```
    Returns:
```

```
        - portfolio_df - the cleaned portfolio dataframe
```

```
    """
```

```
    #deserialize the channels attribute to form each item to a column and their respective values
```

```
    #Convert categorical variable into dummy/indicator variables.
```

```
    #credit to https://stackoverflow.com/questions/70139966/how-to-transform-dataframe-columns-to-dummies
```

```

portfolio_stacked = df.set_index('id')['channels'].map(Counter).apply(pd.Series).fill
#merge the stacked df and the main df to form the collective data
portfolio_df = df.merge(portfolio_stacked,how='left',on='id')
#drop the original channels data
portfolio_df.drop(columns='channels', inplace=True)

#portfolio represent Offers sent during 30-day test period (10 offers x 6 fields)
#we can rename the portfolio id as the offer_id
portfolio_df.rename(columns={'id':'offer_id'}, inplace=True)

#convert the offertypes to columns
offer_types = pd.get_dummies(portfolio_df['offer_type'].str.lower(), prefix="offer")
portfolio_df = pd.concat([portfolio_df,offer_types],axis=1)
portfolio_df.drop(columns='offer_type', inplace=True)

#drop duplicated if any
portfolio_df.drop_duplicates(inplace=True)

return portfolio_df

```

```
portfolio_df = clean_portfolio(portfolio)
```

```
In [7]: portfolio_df.head()
```

```

Out[7]:   difficulty  duration                                offer_id  reward  email  \
0           10         7  ae264e3637204a6fb9bb56bc8210ddfd         10        1
1           10         5  4d5c57ea9a6940dd891ad53e9dbe8da0         10        1
2            0         4  3f207df678b143eea3cee63160fa8bed          0        1
3            5         7  9b98b8c7a33c4b65b9aebfe6a799e6d9          5        1
4           20        10  0b1e1539f2cc45b7b9fa7c272da2e1d7          5        1

   mobile  social  web  offer_bogo  offer_discount  offer_informational
0        1        1    0           1              0                    0
1        1        1    1           1              0                    0
2        1        0    1           0              0                    1
3        1        0    1           1              0                    0
4        0        0    1           0              1                    0

```

```
In [8]: import datetime
```

```
def clean_profile(df):
```

```
'''
```

```
    Description: Clean the profile data frame.
```

- Format the date field to a standard date format
- Create a columns name for all confirmed users using afte inspecting the dataset
- Extract dummy or indicator variable from the categorical attribute gender
- Merge the dummy data with the main dataset
- Drop any duplicates if need be

```
Input:
```

```

    - df: original profile dataframe
Returns:
    - profile_df - the cleaned profile dataframe
'''
df['became_member_on'] = pd.to_datetime(df['became_member_on'], format='%Y%m%d')
df['valid'] = (df.age != 118).astype(int)

# Convert categorical variable line gender into dummy/indicator variables.
#https://pandas.pydata.org/docs/reference/api/pandas.get_dummies.html
dummies = pd.get_dummies(df['gender'].str.lower(), prefix="gender").groupby(df['id'])
#dummies_years = pd.get_dummies(df['became_member_on'].apply(lambda x: x.year)).groupby(df['id'])
df['became_member_on'] = df['became_member_on'].apply(lambda x: int((pd.to_datetime(x).year)))

#merge the stacked df and the main df to form the collective data
profile_df = df.merge(dummies,how='left',on='id')
#profile_df = df.merge(dummies_years,how='left',on='id')

#we can rename the profile id as the user_id
profile_df.rename(columns={'id':'customer_id'}, inplace=True)

#drop duplicated if any
profile_df.drop_duplicates(inplace=True)
#profile_df.dropna(inplace=True)

return profile_df

#Test the clean_profile function
profile_df = clean_profile(profile)
profile_df.head()

```

```

Out[8]:
   age  became_member_on  gender  customer_id  income \
0  118                1872   None  68be06ca386d4c31939f3a4f0e3dd783    NaN
1   55                1719    F  0610b486422d4921ae7d2bf64640c50b  112000.0
2  118                1357   None  38fe809add3b4fcf9315a9694bb96ff5    NaN
3   75                1786    F  78afa995795e4d85b5d9ceeca43f5fef  100000.0
4  118                1699   None  a03223e636434f42ac4c3df47e8bac43    NaN

   valid  gender_f  gender_m  gender_o
0      0         0         0         0
1      1         1         0         0
2      0         0         0         0
3      1         1         0         0
4      0         0         0         0

```

```

In [9]: transcript.head(5)

```

```

Out[9]:
   event  person  time \
0  offer received  78afa995795e4d85b5d9ceeca43f5fef    0

```

1	offer received	a03223e636434f42ac4c3df47e8bac43	0
2	offer received	e2127556f4f64592b11af22de27a7932	0
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0
4	offer received	68617ca6246f4fbc85e91a2a49552598	0

		value
0	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}	
1	{'offer id': '0b1e1539f2cc45b7b9fa7c272da2e1d7'}	
2	{'offer id': '2906b810c7d4411798c6938adc9daaa5'}	
3	{'offer id': 'fafdc668e3743c1bb461111dcafc2a4'}	
4	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}	

```
In [10]: def clean_transcript(df):
    """
    Dataset contains the following attributes
    - person: (string/hash)
    - event: (string) offer received, offer viewed, transaction, offer completed
    - value: (dictionary) different values depending on event type (offers and transaction)
    offer id: (string/hash) not associated with any "transaction"
    amount: (numeric) money spent in "transaction"
    reward: (numeric) money gained from "offer completed"
    (numeric) hours after start of test
    Description: Clean the profile data frame.
        - replace every space with underscore (_) to get a valid column name
        - Extract dummy or indicator variable from the categorical attribute event
        - Merge the dummy data with the main dataset

    Input:
        - df: original transcript dataframe
    Returns:
        - transcript_df - the cleaned transcript dataframe
    """
    # Split event into several dummy columns
    # create expected fields
    transcript_df = df.copy()
    # introduce an underscore as a word separator for the events
    transcript_df.event = transcript_df.event.str.replace(' ', '_')

    # create offer_id, amount and reward fields from the value collections
    transcript_df['offer_id'] = df['value'].apply(lambda x: x['offer_id'] if 'offer_id' in x
    transcript_df['amount'] = df['value'].apply(lambda x: x['amount'] if 'amount' in x
    transcript_df['reward'] = df['value'].apply(lambda x: x['reward'] if 'reward' in x

    # Drop the value field
    transcript_df.drop('value', axis=1, inplace=True)
    transcript_df.rename(columns={'person': 'customer_id'}, inplace=True)

    # Convert transcript time from hours to days
```

```

transcript_df['time'] = df['time'] / 24

#transform the event attribute to create columns for each customer for the offers
transcript_df = transcript_df.groupby(['customer_id', 'offer_id', 'event'])['time'].count()
transcript_df.reset_index(level=[0,1], inplace = True)

#Replace nan values with 0.0
transcript_df.fillna(0.0, inplace = True)

return transcript_df

transcript_df = clean_transcript(transcript)

In [11]: transcript_df.head()

Out[11]:
event      customer_id      offer_id \
0      0009655768c64bdeb2e877511632db8f  2906b810c7d4411798c6938adc9daaa5
1      0009655768c64bdeb2e877511632db8f  3f207df678b143eea3cee63160fa8bed
2      0009655768c64bdeb2e877511632db8f  5a8bc65990b245e5a138643cd4eb9837
3      0009655768c64bdeb2e877511632db8f  f19421c1d4aa40978ebb69ca19b0e20d
4      0009655768c64bdeb2e877511632db8f  fafdc668e3743c1bb461111dcafc2a4

event  offer_completed  offer_received  offer_viewed
0      1.0              1.0              0.0
1      0.0              1.0              1.0
2      0.0              1.0              1.0
3      1.0              1.0              1.0
4      1.0              1.0              1.0

In [12]: def combine_data(portfolio_df, profile_df, transcript_df):
    """ Combine transaction - transcript_df, demographic - profile_df and offer data -
    Input:
    - transcript_df
    - profile_df
    - portfolio_df
    Output:
    - merged_df: merged dataframe
    """
    #First merge the field that have a common join field (transcript_df and portfolio_df)
    member_transactions = pd.merge(transcript_df, portfolio_df, how='left', on='offer_id')

    #Merge the third dataframe with the common dataset
    merged_df = pd.merge(member_transactions, profile_df, how='left', on='customer_id')

    #Codify the offer_ids
    offer_idx = list(set(merged_df['offer_id']))

    for i in range(len(offer_idx)):

```



```

merged_df['offer_id'] = merged_df['offer_id'].apply(lambda x: '{}'.format(i+1))

#consider an offer to be successful if a viewed offer is completed
merged_df.insert(loc=2, column='offer_successful', value=merged_df['offer_completed'])
merged_df['offer_successful'] = merged_df['offer_successful'].apply(lambda x: 1.0 if x else 0.0)
merged_df.drop(columns=['offer_completed', 'offer_received', 'offer_viewed'], inplace=True)

return merged_df

merged_df = combine_data(portfolio_df, profile_df, transcript_df)

```

In [13]: merged_df.shape

Out[13]: (63288, 21)

4 Descriptive Analysis and Data Exploration

4.0.1 Descriptive Analysis

4.0.2 Lets use graphs, charts and decriptive statistics to get more insight from each data groups

```

In [33]: profile_df_copy = profile_df.copy()
profile_df_copy = profile_df_copy[profile_df_copy['age'] != 118]
profile_df_copy = profile_df_copy[profile_df_copy['income'].isnull() == False]
profile_df_copy.describe()

```

```

Out[33]:


```

	age	became_member_on	income	valid	gender_f \
count	14825.000000	14825.000000	14825.000000	14825.0	14825.000000
mean	54.393524	1865.478988	65404.991568	1.0	0.413423
std	17.383705	419.205158	21598.299410	0.0	0.492464
min	18.000000	1343.000000	30000.000000	1.0	0.000000
25%	42.000000	1551.000000	49000.000000	1.0	0.000000
50%	55.000000	1701.000000	64000.000000	1.0	0.000000
75%	66.000000	2140.000000	80000.000000	1.0	1.000000
max	101.000000	3166.000000	120000.000000	1.0	1.000000

	gender_m	gender_o
count	14825.000000	14825.000000
mean	0.572277	0.014300
std	0.494765	0.118729
min	0.000000	0.000000
25%	0.000000	0.000000
50%	1.000000	0.000000
75%	1.000000	0.000000
max	1.000000	1.000000

From the above descriptive statistics, the average income (revenue generated through the offer) is about USD65,404.99 from a total of 14,825 participants/customers


```
In [20]: plt.hist(profile_df_copy['age'], bins=5);
```



```
In [22]: valid_income = profile_df_copy[profile_df_copy['income'].isna()==False]
```

```
In [23]: plt.hist(valid_income['income'], bins=5)
```

```
Out[23]: (array([ 3438.,  4444.,  3780.,  2205.,   958.]),  
          array([ 30000.,  48000.,  66000.,  84000., 102000., 120000.]),  
          <a list of 5 Patch objects>)
```



```
In [24]: transcript_df.describe()
```

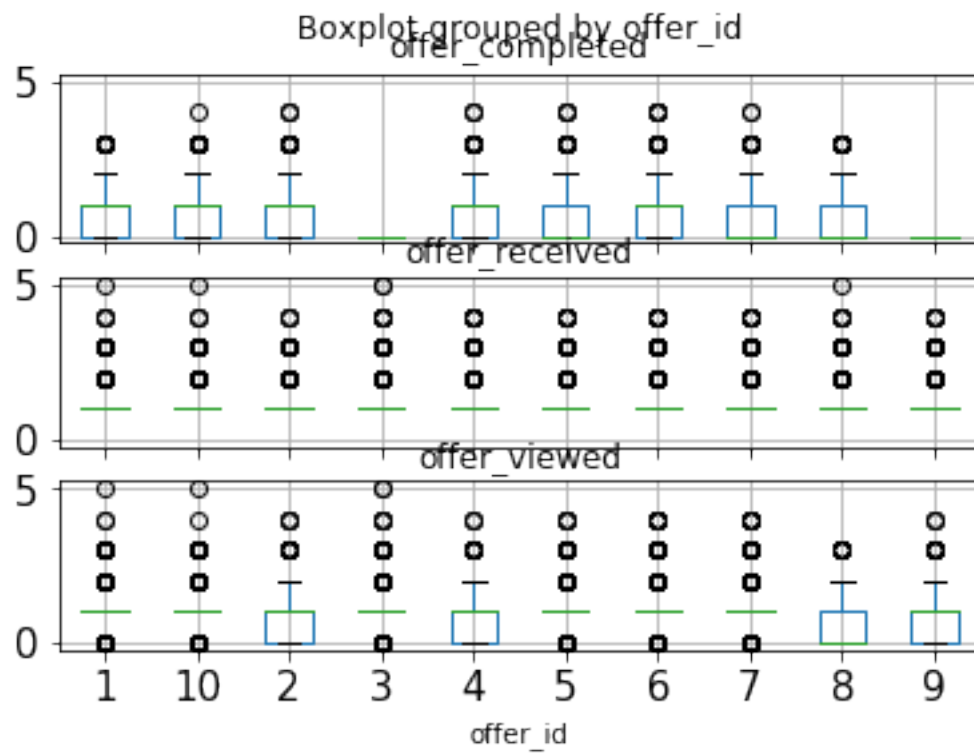
```
Out[24]:
```

	event	offer_completed	offer_received	offer_viewed
count		63288.000000	63288.000000	63288.000000
mean		0.530575	1.205236	0.91210
std		0.637582	0.453699	0.61458
min		0.000000	1.000000	0.000000
25%		0.000000	1.000000	1.000000
50%		0.000000	1.000000	1.000000
75%		1.000000	1.000000	1.000000
max		4.000000	5.000000	5.000000


4.0.3 From the transaction descriptive statistics above, from the mean (offer completed = 0.53 and offer viewed = 0.91, about 39% difference), some of the customers that viewed the received offers were not able to complete the purchase whereas 67% of the offers received were not completed.

```
In [25]: transcript_df_copy = transcript_df.copy()
         offer_idx = list(set(transcript_df_copy.offer_id))
         for i in range(len(offer_idx)):
             transcript_df_copy['offer_id'] = transcript_df_copy['offer_id'].apply(lambda x: '{}')

         boxplot = transcript_df_copy.boxplot(column=['offer_completed', 'offer_received', 'offer_viewed'])
```



```
In [26]: axes = transcript_df_copy.plot.kde()
```



From the above visual, although all the participants received the information to participate in the offers, only about 38% of the population completed their offers.

4.0.4 The merged data show the summary of the three datasets and also reveals the number of records available

In [39]: merged_df.head(10)

```
Out[39]:
```

	customer_id	offer_id	offer_successful	difficulty	\
0	0009655768c64bdeb2e877511632db8f	4	0.0	10	
1	0009655768c64bdeb2e877511632db8f	9	0.0	0	
2	0009655768c64bdeb2e877511632db8f	3	0.0	0	
3	0009655768c64bdeb2e877511632db8f	1	1.0	5	
4	0009655768c64bdeb2e877511632db8f	10	1.0	10	
5	00116118485d4dfda04fdbaba9a87b5c	1	0.0	5	
6	0011e0d4e6b944f998e987f904e8c1e5	8	1.0	20	
7	0011e0d4e6b944f998e987f904e8c1e5	6	1.0	7	
8	0011e0d4e6b944f998e987f904e8c1e5	9	0.0	0	
9	0011e0d4e6b944f998e987f904e8c1e5	3	0.0	0	

	duration	reward	email	mobile	social	web	...	offer_discount	\
0	7	2	1	1	0	1	...	1	
1	4	0	1	1	0	1	...	0	
2	3	0	1	1	1	0	...	0	
3	5	5	1	1	1	1	...	0	
4	10	2	1	1	1	1	...	1	
5	5	5	1	1	1	1	...	0	
6	10	5	1	0	0	1	...	1	
7	7	3	1	1	1	1	...	1	
8	4	0	1	1	0	1	...	0	
9	3	0	1	1	1	0	...	0	

	offer_informational	age	became_member_on	gender	income	valid	\		
0		0	33		1804	M	72000.0	1	
1		1	33		1804	M	72000.0	1	
2		1	33		1804	M	72000.0	1	
3		0	33		1804	M	72000.0	1	
4		0	33		1804	M	72000.0	1	
5		0	118		1435	None	NaN	0	
6		0	40		1541	0	57000.0	1	
7		0	40		1541	0	57000.0	1	
8		1	40		1541	0	57000.0	1	
9		1	40		1541	0	57000.0	1	

	gender_f	gender_m	gender_o
0	0	1	0
1	0	1	0
2	0	1	0

3	0	1	0
4	0	1	0
5	0	0	0
6	0	0	1
7	0	0	1
8	0	0	1
9	0	0	1

[10 rows x 21 columns]

In [29]: merged_df.describe()

```
Out[29]:
```

	offer_successful	difficulty	duration	reward	email \
count	63288.000000	63288.000000	63288.000000	63288.000000	63288.0
mean	0.386487	7.711572	6.504819	4.206232	1.0
std	0.486948	5.541480	2.203565	3.402914	0.0
min	0.000000	0.000000	3.000000	0.000000	1.0
25%	0.000000	5.000000	5.000000	2.000000	1.0
50%	0.000000	10.000000	7.000000	5.000000	1.0
75%	1.000000	10.000000	7.000000	5.000000	1.0
max	1.000000	20.000000	10.000000	10.000000	1.0

	mobile	social	web	offer_bogo	offer_discount \
count	63288.000000	63288.000000	63288.000000	63288.000000	63288.000000
mean	0.899286	0.599529	0.799425	0.400092	0.400013
std	0.300952	0.489998	0.400434	0.489921	0.489904
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	1.000000	0.000000	0.000000
50%	1.000000	1.000000	1.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	offer_informational	age	became_member_on	income \
count	63288.000000	63288.000000	63288.000000	55222.000000
mean	0.199896	62.462110	1859.721053	65388.595125
std	0.399925	26.729957	411.659882	21626.373809
min	0.000000	18.000000	1343.000000	30000.000000
25%	0.000000	45.000000	1551.000000	49000.000000
50%	0.000000	58.000000	1701.000000	63000.000000
75%	0.000000	73.000000	2134.000000	80000.000000
max	1.000000	118.000000	3166.000000	120000.000000

	valid	gender_f	gender_m	gender_o
count	63288.000000	63288.000000	63288.000000	63288.000000
mean	0.872551	0.360384	0.500016	0.012151
std	0.333478	0.480116	0.500004	0.109560
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000

50%	1.000000	0.000000	1.000000	0.000000
75%	1.000000	1.000000	1.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

5 Determine which demographic groups respond best to which offer type

```
In [30]: merged_df[merged_df['offer_successful'] == 1].offer_id.value_counts()
```

```
Out[30]: 10    4433
         6    4313
         1    3655
         5    2843
         7    2825
         2    2504
         4    2415
         8    1472
         Name: offer_id, dtype: int64
```

Provide descriptive statistics for all cases where customer's income was not reflected

```
In [17]: merged_df[merged_df['income'].isnull()][['age', 'income']].describe()
```

```
Out[17]:
```

	age	income
count	8066.0	0.0
mean	118.0	NaN
std	0.0	NaN
min	118.0	NaN
25%	118.0	NaN
50%	118.0	NaN
75%	118.0	NaN
max	118.0	NaN

Provide descriptive statistics for all cases where customers' income exist

```
In [31]: merged_df[merged_df['income'].notnull()][['age', 'income']].describe()
```

```
Out[31]:
```

	age	income
count	55222.000000	55222.000000
mean	54.349969	65388.595125
std	17.392733	21626.373809
min	18.000000	30000.000000
25%	42.000000	49000.000000
50%	55.000000	63000.000000
75%	66.000000	80000.000000
max	101.000000	120000.000000


Evaluate year membership record by year using the field `became_member_on`

```
In [35]: #start_year = merged_df['became_member_on'].apply(lambda x: x.year).value_counts()  
        #start_year
```

5.0.1 Data Visualization (contd)


```
In [36]: merged_df.plot.scatter(x='age', y='income', c='offer_successful', cmap='Blues')
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7efcd17d8128>
```




```
In [37]: sns.distplot(merged_df['difficulty'].dropna())
```

```
Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x7efcd18982b0>
```


```
In [38]: #get the summary of ages of the participants
age_df = merged_df.groupby(['gender', 'age']).size().sort_values(ascending=False) .reset_index()
plt.figure(figsize=(15, 5))
sns.barplot(x='age', y='count', hue='gender', data=age_df)
plt.xlabel('Member Age')
plt.ylabel('Count')
plt.xticks(rotation=90)

Out[38]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67,
        68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83]),
<a list of 84 Text xticklabel objects>)
```



5.0.2 Plot distribution of age by gender to see the response of each gender to offers accross ages


```
In [21]: sns.distplot(merged_df[merged_df['gender']=='M']['age'],hist=False,color="g", kde_kws={
sns.distplot(merged_df[merged_df['gender']=='F']['age'],hist=False,color="b", kde_kws={
sns.distplot(merged_df[merged_df['gender']=='O']['age'],hist=False,color="r", kde_kws={
plt.title('Distribution of Age by Gender')
plt.gca().get_yaxis().set_visible(False)
plt.legend(['Male', 'Female', 'Others'],frameon=False);
```



From the diagram above, maximum patronage occurred at age between 40-70 accross all gender

Plot distribution of income by gender to see the response of each gender to offers based on income

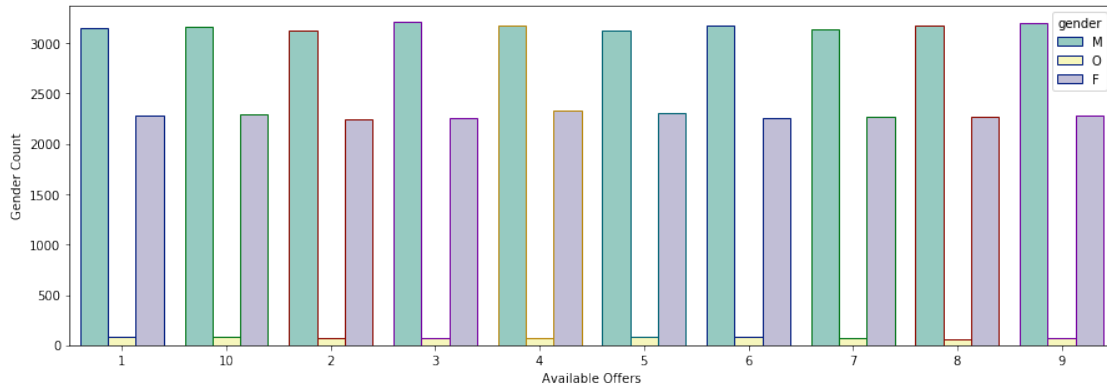
```
In [22]: sns.distplot(merged_df[merged_df['gender']=='M']['income'],hist=False,color="g", kde_kws={
sns.distplot(merged_df[merged_df['gender']=='F']['income'],hist=False,color="b", kde_kws={
sns.distplot(merged_df[merged_df['gender']=='O']['income'],hist=False,color="r", kde_kws={
plt.title('Distribution of Income by Gender')
plt.gca().get_yaxis().set_visible(False)
plt.legend(['Male','Female','Others'],frameon=False);
```



5.0.3 Count of offers on Gender basis

```
In [23]: #gender_df = merged_df[merged_df['gender']=='M' or merged_df['gender']=='F']
plt.figure(figsize=(15, 5))
ax = sns.countplot(x="offer_id", hue="gender", data=merged_df, palette="Set3", edgecolor=
plt.xlabel('Available Offers')
plt.ylabel('Gender Count')
```

```
Out[23]: Text(0,0.5,'Gender Count')
```




5.0.4 Analysis of successful and unsuccessful offers based on status, income, duration, difficulty and reward

1 Based on Status

```
In [24]: plt.figure(figsize=(15, 5))
sns.countplot(x='offer_id', hue='offer_successful', data=merged_df, palette="Set3", edgecolor='green')
plt.legend(['Successful', 'Unsuccessful'], frameon=False)
plt.title('Count of Offer Type')
```

```
Out[24]: Text(0.5,1,'Count of Offer Type')
```



2. Based on Income


```
In [25]: sns.distplot(merged_df[merged_df['offer_successful'] == 1]['income'], hist=False, color="blue")
sns.distplot(merged_df[merged_df['offer_successful'] == 0]['income'], hist=False, color="yellow")
plt.legend(['Successful', 'Unsuccessful'], frameon=False)
plt.gca().get_yaxis().set_visible(False)
plt.title('Distribution of Status of Offers by Income')
```

```

/opt/conda/lib/python3.6/site-packages/statsmodels/nonparametric/kde.py:454: RuntimeWarning: inv
  X = X[np.logical_and(X>clip[0], X<clip[1])] # won't work for two columns.
/opt/conda/lib/python3.6/site-packages/statsmodels/nonparametric/kde.py:454: RuntimeWarning: inv
  X = X[np.logical_and(X>clip[0], X<clip[1])] # won't work for two columns.

```

Out[25]: Text(0.5,1,'Distribution of Status of Offers by Income')



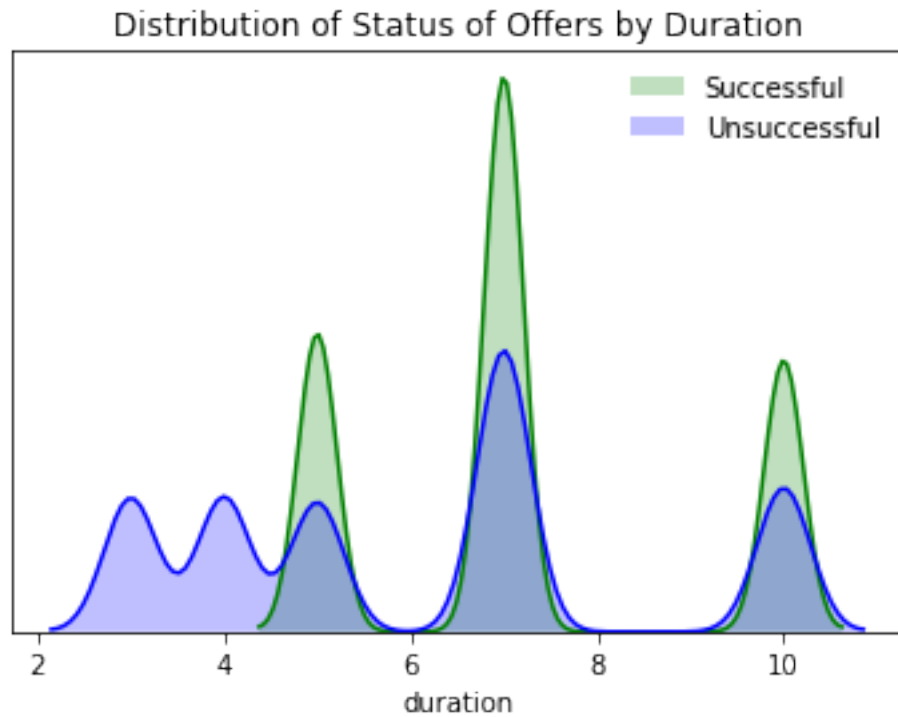
3. Based on Duration

```

In [26]: sns.distplot(merged_df[merged_df['offer_successful'] > 0]['duration'],hist=False,color=
sns.distplot(merged_df[merged_df['offer_successful']<= 0]['duration'],hist=False,color=
plt.legend(['Successful', 'Unsuccessful'],frameon=False)
plt.gca().get_yaxis().set_visible(False)
plt.title('Distribution of Status of Offers by Duration')

```


Out[26]: Text(0.5,1,'Distribution of Status of Offers by Duration')



4. Based on difficulty

```
In [27]: sns.distplot(merged_df[merged_df['offer_successful'] > 0]['difficulty'], hist=False, color='green')
sns.distplot(merged_df[merged_df['offer_successful'] <= 0]['difficulty'], hist=False, color='blue')
plt.legend(['Successful', 'Unsuccessful'], frameon=False)
plt.gca().get_yaxis().set_visible(False)
plt.title('Distribution of Status of Offers by difficulty')
```


```
Out[27]: Text(0.5, 1, 'Distribution of Status of Offers by difficulty')
```



5. Based on reward

```
In [28]: sns.distplot(merged_df[merged_df['offer_successful'] > 0]['reward'], hist=False, color="g")
sns.distplot(merged_df[merged_df['offer_successful'] <= 0]['reward'], hist=False, color="b")
plt.legend(['Successful', 'Unsuccessful'], frameon=False)
plt.gca().get_yaxis().set_visible(False)
plt.title('Distribution of Status of Offers by reward')
```

```
Out[28]: Text(0.5,1,'Distribution of Status of Offers by reward')
```



5.0.5 Modeling

Let us explore the some modelling techniques to perform some exploratory analysis. Modeling techniques are used in data evaluation and analysis. We would reveal some important findings the best performing offers, etc using these modelling techniques.

```
In [40]: merged_df.loc[merged_df['offer_successful'] == 1].head()
```

```
Out[40]:
```

	customer_id	offer_id	offer_successful	difficulty	\
3	0009655768c64bdeb2e877511632db8f	1	1.0	5	
4	0009655768c64bdeb2e877511632db8f	10	1.0	10	
6	0011e0d4e6b944f998e987f904e8c1e5	8	1.0	20	
7	0011e0d4e6b944f998e987f904e8c1e5	6	1.0	7	
10	0011e0d4e6b944f998e987f904e8c1e5	2	1.0	5	

	duration	reward	email	mobile	social	web	...	offer_discount	\
3	5	5	1	1	1	1	...	0	
4	10	2	1	1	1	1	...	1	
6	10	5	1	0	0	1	...	1	
7	7	3	1	1	1	1	...	1	
10	7	5	1	1	0	1	...	0	

	offer_informational	age	became_member_on	gender	income	valid	\
3	0	33	1804	M	72000.0	1	

4	0	33	1804	M	72000.0	1
6	0	40	1541	0	57000.0	1
7	0	40	1541	0	57000.0	1
10	0	40	1541	0	57000.0	1

	gender_f	gender_m	gender_o
3	0	1	0
4	0	1	0
6	0	0	1
7	0	0	1
10	0	0	1

[5 rows x 21 columns]

In [41]: merged_df.drop(columns=['became_member_on', 'gender', 'valid', 'offer_successful']).iloc

Out[41]:

	offer_id	difficulty	duration	reward	email	mobile	social	web	\
0	4	10	7	2	1	1	0	1	
1	9	0	4	0	1	1	0	1	
2	3	0	3	0	1	1	1	0	
3	1	5	5	5	1	1	1	1	
4	10	10	10	2	1	1	1	1	

	offer_bogo	offer_discount	offer_informational	age	income	gender_f	\
0	0	1	0	33	72000.0	0	
1	0	0	1	33	72000.0	0	
2	0	0	1	33	72000.0	0	
3	1	0	0	33	72000.0	0	
4	0	1	0	33	72000.0	0	

	gender_m	gender_o
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

In [42]: merged_df.iloc[:,1].head()
merged_df.iloc[:,10:].head()

Out[42]:

	offer_bogo	offer_discount	offer_informational	age	became_member_on	\
0	0	1	0	33	1804	
1	0	0	1	33	1804	
2	0	0	1	33	1804	
3	1	0	0	33	1804	
4	0	1	0	33	1804	

	gender	income	valid	gender_f	gender_m	gender_o
0	M	72000.0	1	0	1	0

1	M	72000.0	1	0	1	0
2	M	72000.0	1	0	1	0
3	M	72000.0	1	0	1	0
4	M	72000.0	1	0	1	0

```
In [43]: merged_df.isnull().sum()
```

```
Out[43]: customer_id      0
         offer_id        0
         offer_successful  0
         difficulty      0
         duration        0
         reward          0
         email           0
         mobile          0
         social          0
         web             0
         offer_bogo      0
         offer_discount  0
         offer_informational 0
         age             0
         became_member_on 0
         gender          8066
         income          8066
         valid           0
         gender_f        0
         gender_m        0
         gender_o        0
         dtype: int64
```

5.0.6 Create a model function that would remove the gender field since dummy values had earlier been generated for the field, then fill missing values in the income field with the income mean value

```
In [50]: def create_model_data(df):
         '''
         Description: Create a function that would return the return X variables and the pr
         Input
         - merged_df
         output
         - two dataframes X and Y
             X - return the collection of the merged data.
             Y - return the unique offer types.
         '''

         #drop categorical culumns from the experiment dataset
         my_model = df.copy()
         model_df = my_model.drop(columns=['gender', 'valid'])
         #We want to look at only successful offers
```

```

model_df = model_df.loc[model_df['offer_successful'] == 1]
#fill the missing values with mean of the field (data cleaning and working with miss
model_df['income'] = model_df['income'].fillna(model_df['income'].mean())
#gat all th rows from the 3rd colmun, focus on the quantitative data
X = model_df.iloc[:,3:]

#The the labels of all Offers
Y= model_df.offer_id

return X, Y

```

```

In [51]: features, target_labels = create_model_data(merged_df)
         features.shape, target_labels.shape

```

```

Out[51]: ((24460, 16), (24460,))

```

```

In [52]: target_labels.head()

```

```

Out[52]: 3      1
         4     10
         6      8
         7      6
        10      2
         Name: offer_id, dtype: object

```

```

In [53]: features.head()

```

```

Out[53]:
   difficulty  duration  reward  email  mobile  social  web  offer_bogo \
3           5         5       5      1      1      1      1          1
4          10        10       2      1      1      1      1          0
6          20        10       5      1      0      0      1          0
7           7         7       3      1      1      1      1          0
10          5         7       5      1      1      0      1          1

   offer_discount  offer_informational  age  became_member_on  income \
3                0                   0  33                1804  72000.0
4                1                   0  33                1804  72000.0
6                1                   0  40                1541  57000.0
7                1                   0  40                1541  57000.0
10               0                   0  40                1541  57000.0

   gender_f  gender_m  gender_o
3          0         1         0
4          0         1         0
6          0         0         1
7          0         0         1
10         0         0         1

```

5.0.7 1. Using Random Forrest Classifier:

In using this algorithm, it allows us to recommend the best offers for other customers using divide-and-conquer approach. It is a highly accurate and robust method because of the number of decision trees participating in the process . It does not suffer from the overfitting problem, takes the average of all the predictions, and cancels out the biases.

It works in four steps:

- Select random samples from a given dataset.
- Construct a decision tree for each sample and get a prediction result from each decision tree.
- Perform a vote for each predicted result.
- Select the prediction result with the most votes as the final prediction.

In []:

In [54]: *# Estimate the Scaling our inputs*

```
min_max_scaler = preprocessing.MinMaxScaler()  
X = min_max_scaler.fit_transform(features)
```

In [55]: *#Use the experiment data to estimate the training and the test data sets*

```
X_train, X_test, y_train, y_test = \  
train_test_split(features, target_labels, test_size=.25, random_state=21)
```

```
X_train.shape, X_test.shape
```

Out[55]: ((18345, 16), (6115, 16))

In [56]: *# Scaling our inputs*

```
min_max_scaler2 = preprocessing.MinMaxScaler()  
X_train_fit = min_max_scaler2.fit_transform(X_train)  
X_test_fit = min_max_scaler2.fit_transform(X_test)
```

In [57]: *#Create a Gaussian Classifier*

```
clf = RandomForestClassifier(n_estimators=100)  
clf.fit(X_train_fit, y_train)
```

Out[57]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)

In [58]: *#Predict possible Offers using the test dataset*

```
y_pred=clf.predict(X_test_fit)
```

In [61]: y_pred

Out[61]: array(['2', '6', '10', ..., '6', '6', '1'], dtype=object)

5.0.8 Use Confusion Matrix to measure the accuracy of the prediction

- The Confusion Matrix allows us to visualize the performance of the classification machine learning models. With this visualization technique, we can get a better idea of how our machine learning model is performing
- We can look at a confusion matrix to see how much accurate classification and misclassification there was. The diagonal indicates correct % of classification. Off diagonals indicate the % of times the model misclassified an offer.

```
In [62]: #get the confusion matrix and the accuracy score
confusion = confusion_matrix(y_test,y_pred)
accuracy_score(y_test,y_pred)
```

```
Out[62]: 1.0
```

```
In [63]: confusion
```


```
Out[63]: array([[ 921,    0,    0,    0,    0,    0,    0,    0],
 [    0, 1145,    0,    0,    0,    0,    0,    0],
 [    0,    0,  612,    0,    0,    0,    0,    0],
 [    0,    0,    0,  607,    0,    0,    0,    0],
 [    0,    0,    0,    0,  691,    0,    0,    0],
 [    0,    0,    0,    0,    0, 1070,    0,    0],
 [    0,    0,    0,    0,    0,    0,  711,    0],
 [    0,    0,    0,    0,    0,    0,    0,  358]])
```

```
In [64]: ax = sns.heatmap(confusion/np.sum(confusion), annot=True, fmt='.1%', cmap='Blues')
#ax = sns.heatmap(confusion, cmap='Blues')
```

```
ax.set_title('Offers Prediction Confusion Matrix with labels\n\n');
ax.set_xlabel('\nPredicted Offers')
ax.set_ylabel('Actual Offers ');
plt.figure(figsize = (25,20))
```

```
## Display the visualization of the Confusion Matrix.
plt.show()
```

Offers Prediction Confusion Matrix with labels



<matplotlib.figure.Figure at 0x7efcec9f35f8>

5.1 Using Model to classify individual gata groups

5.1.1 1. To determine which demographic group that responded best to each offer

- Generate model data using the demographic collection

In [66]: X,Y = create_model_data(merged_df)

```
#Only keep inforamtion in the X dataframe that refers to the user.
X = X.iloc[:, 10:]
X.head()
```

```
Out[66]:
```

	age	became_member_on	income	gender_f	gender_m	gender_o
3	33	1804	72000.0	0	1	0
4	33	1804	72000.0	0	1	0
6	40	1541	57000.0	0	0	1
7	40	1541	57000.0	0	0	1
10	40	1541	57000.0	0	0	1

```
In [67]: #Generate the train and test data for prediction
X_train, X_test, y_train, y_test = train_test_split(\
    X,Y,test_size=.2,random_state=21)

#Normalize the data
sc = preprocessing .StandardScaler()
X_train=sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)

#Instansiate Classifier
classifier = RandomForestClassifier(n_estimators=100)

#Train Classifier
classifier.fit(X_train,y_train)

Out[67]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
    max_depth=None, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
    oob_score=False, random_state=None, verbose=0,
    warm_start=False)
```

```
In [68]: y_pred = classifier.predict(X_test)

print('Model accuracy: {0:0.2f}'.format(accuracy_score(y_test,y_pred)))
print(classification_report(y_test,y_pred))
```

Model accuracy: 0.06

	precision	recall	f1-score	support
1	0.06	0.05	0.05	737
10	0.12	0.12	0.12	909
2	0.04	0.03	0.03	485
4	0.02	0.02	0.02	487
5	0.02	0.02	0.02	554
6	0.10	0.11	0.11	864
7	0.03	0.03	0.03	571
8	0.00	0.00	0.00	285
avg / total	0.06	0.06	0.06	4892


```
In [51]: #Generate the confusion matrix
cf_matrix_std = confusion_matrix(y_test, y_pred)
```

```
In [52]: cf_matrix_std
```

```
Out[52]: array([[ 12,  62, 105,  33, 105,  54,  61,  53],
 [ 60,  47, 171,  39, 176,  62,  88,  94],
 [ 82, 166, 112,  35, 235,  75, 101, 103],
 [ 41,  44,  34,   2,  53,  30,  39,  42],
 [103, 155, 193,  41,  97,  82, 100,  93],
 [ 43,  74, 121,  22, 101,  12,  50,  64],
 [ 60,  85, 116,  29, 118,  59,  21,  83],
 [ 46, 112, 103,  30, 118,  43,  91,  11]])
```

```
In [53]: ax = sns.heatmap(cf_matrix_std/np.sum(cf_matrix_std), annot=True, fmt='.1%', cmap='Blue')
ax.set_title('Confusion Matrix with labels based on Demography\n\n');
ax.set_xlabel('\nPredicted Offers')
ax.set_ylabel('Actual Offers ');
## Ticket labels - List must be in alphabetical order
ax.xaxis.set_ticklabels(y_pred)
ax.yaxis.set_ticklabels(y_pred)
## Display the visualization of the Confusion Matrix.
plt.show()
```

Confusion Matrix with labels based on Demography



5.1.2 To determine which Offer performed best among the offer types

```
In [54]: X,Y = create_model_data(merged_df)
```



```
In [59]: X = X.iloc[:, :10]
        X.head()
```

```
Out[59]:
```

	difficulty	duration	reward	email	mobile	social	web	offer_bogo	\
3	5	5	5	1	1	1	1	1	
4	10	10	2	1	1	1	1	0	
6	20	10	5	1	0	0	1	0	
7	7	7	3	1	1	1	1	0	
10	5	7	5	1	1	0	1	1	

	offer_discount	offer_informational
3	0	0
4	1	0
6	1	0
7	1	0
10	0	0

```
In [60]: #Generate the train and test data for prediction
        X_train, X_test, y_train, y_test = train_test_split(X,Y,test_size=.2,random_state=21)

        #Normalize the data
        sc = preprocessing.StandardScaler()
        X_train=sc.fit_transform(X_train)
        X_test = sc.fit_transform(X_test)

        #Instansiate Classifier
        classifier = RandomForestClassifier(n_estimators=20,criterion='entropy',random_state=42)

        #Train Classifier
        classifier.fit(X_train,y_train)
```

```
Out[60]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=20, n_jobs=1,
                                oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [61]: y_pred = classifier.predict(X_test)

        print('Model accuracy: {0:0.2f}'.format(accuracy_score(y_test,y_pred)))
        print(classification_report(y_test,y_pred))
```

Model accuracy: 1.00

	precision	recall	f1-score	support
10	1.00	1.00	1.00	485
2	1.00	1.00	1.00	737

3	1.00	1.00	1.00	909
4	1.00	1.00	1.00	285
6	1.00	1.00	1.00	864
7	1.00	1.00	1.00	487
8	1.00	1.00	1.00	571
9	1.00	1.00	1.00	554
avg / total	1.00	1.00	1.00	4892


```
In [62]: #Generate the confusion matrix
         cf_matrix_std_offers = confusion_matrix(y_test, y_pred)
```

```
In [63]: cf_matrix_std_offers
```

```
Out[63]: array([[485,  0,  0,  0,  0,  0,  0,  0],
                [ 0, 737,  0,  0,  0,  0,  0,  0],
                [ 0,  0, 909,  0,  0,  0,  0,  0],
                [ 0,  0,  0, 285,  0,  0,  0,  0],
                [ 0,  0,  0,  0, 864,  0,  0,  0],
                [ 0,  0,  0,  0,  0, 487,  0,  0],
                [ 0,  0,  0,  0,  0,  0, 571,  0],
                [ 0,  0,  0,  0,  0,  0,  0, 554]])
```

```
In [64]: ax = sns.heatmap(cf_matrix_std_offers/np.sum(cf_matrix_std_offers), annot=True, fmt='.1f',
                        ax.set_title('Confusion Matrix with labels based on Offers\n\n');
                        ax.set_xlabel('\nPredicted Offers')
                        ax.set_ylabel('Actual Offers ');
                        ## Ticket labels - List must be in alphabetical order
                        ax.xaxis.set_ticklabels(y_pred)
                        ax.yaxis.set_ticklabels(y_pred)
                        ## Display the visualization of the Confusion Matrix.
                        plt.show()
```

Confusion Matrix with labels based on Offers



5.1.3 Using Logistic Regression

- Gradient Boosting uses an ensemble of decision trees to predict a target label. However, unlike AdaBoost, the Gradient Boost trees have a depth larger than 1
- Calculate the average of the target label
- Calculate the residuals: For every sample, we calculate the residual with the proceeding formula.
- $\text{residual} = \text{actual value} - \text{predicted value}$
- Construct a decision tree
- Predict the target label using all of the trees within the ensemble
- Once trained, use all of the trees in the ensemble to make a final prediction as to the value of the target variable

5.1.4 Use GridSearchCV to discover optimal parameters

<https://www.kaggle.com/enespolat/grid-search-with-logistic-regression> https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html

```
In [69]: #create a model_pipeline
def model_pipeline(X_features, y_target):
    ...
```

```

inputs:
- X_features & y_target dataframe

outputs:
- Splits features and target dataframe to train and test sets, performs feature scaling
- Outputs X_train, X_test, y_train and y_test dataframes
'''

#split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_features,y_target, test_size=0.2)

#fit and transform scaling on training data
scaler= preprocessing.StandardScaler()
X_train=scaler.fit_transform(X_train)

#scale test data
X_test=scaler.transform(X_test)

return X_train,X_test,y_train, y_test

```

```

In [70]: #define Grid Search function
def grid_search_random_forest(X,y):
    '''
    input:
    - X,y: training datasets for X and y
    output:
    - dictionary with best parameters for random forest model
    '''

    params={'max_features': ['auto', 'sqrt'],
            'max_depth' : [5,10,15,20],
            'n_estimators': [20,30,40,50],
            'min_samples_split': [2, 10, 20],
            'min_samples_leaf': [2, 10,15, 20],
            }

    gs_cv = GridSearchCV(RandomForestClassifier(random_state=2), params)
    gs_cv.fit(X, y)
    gs_cv.best_params_

    return gs_cv.best_params_

```

```

In [77]: # Split the dataset into training and test dataset
X,Y = create_model_data(merged_df)
x_train, x_test, y_train, y_test = train_test_split(X, Y, random_state=1)

# Create a Logistic Regression Object, perform Logistic Regression
#using some information from
#https://www.kaggle.com/code/enespolat/grid-search-with-logistic-regression/notebook

```

#<https://www.datarmatics.com/data-science/how-to-perform-logistic-regression-in-pythons>

```
params = {
    'penalty': ['l1', 'l2'],
    'C': [1, 10, 100, 1000]
}

grid_cv = GridSearchCV(LogisticRegression(C=1.0, class_weight=None, dual=False, fit_int
                        intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                        verbose=3, warm_start=False), params, verbose=3, n_jobs=-1, cv=3)
```

```
In [78]: grid_cv.fit(x_train, y_train)
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

[CV] C=1, penalty=l1 ...

[LibLinear][CV] ... C=1, penalty=l1, score=1.0, total= 0.9s

[CV] C=1, penalty=l1 ...

[LibLinear]

[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.9s remaining: 0.0s

[CV] ... C=1, penalty=l1, score=1.0, total= 0.9s

[CV] C=1, penalty=l1 ...

[LibLinear]

[Parallel(n_jobs=-1)]: Done 2 out of 2 | elapsed: 1.8s remaining: 0.0s

[CV] ... C=1, penalty=l1, score=1.0, total= 0.9s

[CV] C=1, penalty=l2 ...

[LibLinear][CV] ... C=1, penalty=l2, score=0.9344557044785877, total= 1.2s

[CV] C=1, penalty=l2 ...

[LibLinear][CV] ... C=1, penalty=l2, score=0.8220477592410861, total= 1.2s

[CV] C=1, penalty=l2 ...

[LibLinear][CV] ... C=1, penalty=l2, score=0.891542614101096, total= 1.2s

[CV] C=10, penalty=l1 ...

[LibLinear][CV] ... C=10, penalty=l1, score=1.0, total= 0.6s

[CV] C=10, penalty=l1 ...

[LibLinear][CV] ... C=10, penalty=l1, score=1.0, total= 0.6s

[CV] C=10, penalty=l1 ...

[LibLinear][CV] ... C=10, penalty=l1, score=1.0, total= 0.7s

[CV] C=10, penalty=l2 ...

[LibLinear][CV] ... C=10, penalty=l2, score=0.9382151029748284, total= 1.2s

[CV] C=10, penalty=l2 ...

[LibLinear][CV] ... C=10, penalty=l2, score=0.8544324501144913, total= 1.3s

[CV] C=10, penalty=l2 ...

[LibLinear][CV] ... C=10, penalty=l2, score=0.8378864714542777, total= 1.1s

```

[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=1.0, total= 0.7s
[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=1.0, total= 0.6s
[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=1.0, total= 0.5s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.9488394900294214, total= 1.2s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.8609748119070985, total= 1.3s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.8885980696875512, total= 1.1s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=1.0, total= 0.6s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=1.0, total= 0.6s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=1.0, total= 0.7s
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.940339980385747, total= 1.2s
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.8132155708210664, total= 1.2s
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.8845084246687388, total= 1.1s
[LibLinear]

```

```

[Parallel(n_jobs=-1)]: Done 24 out of 24 | elapsed: 23.0s finished

```

```

Out[78]: GridSearchCV(cv=3, error_score='raise',
                      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                                                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                                                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                                                    verbose=3, warm_start=False),
                      fit_params=None, iid=True, n_jobs=-1,
                      param_grid={'penalty': ['l1', 'l2'], 'C': [1, 10, 100, 1000]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=3)

```

```

In [79]: grid_cv.best_params_

```

```

Out[79]: {'C': 1, 'penalty': 'l1'}

```

```

In [80]: logistic_y_pred = grid_cv.predict(x_test)
         logistic_accuracy = accuracy_score(y_test, logistic_y_pred)
         #logistic_f1 = f1_score(y_test, logistic_y_pred)

```

```

In [81]: accuracy_score(y_test, logistic_y_pred)

```

```

Out[81]: 1.0

```

```
In [82]: #y_pred = log_model.predict(X_test)
conf_matrix = confusion_matrix(y_test, logistic_y_pred)
print(classification_report(y_test, logistic_y_pred))
print(confusion_matrix(y_test, logistic_y_pred))
print(accuracy_score(y_test, logistic_y_pred))

print("tuned hpyerparameters :(best parameters) ", grid_cv.best_params_)
print("accuracy :",grid_cv.best_score_)
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	919
10	1.00	1.00	1.00	1098
2	1.00	1.00	1.00	610
4	1.00	1.00	1.00	603
5	1.00	1.00	1.00	713
6	1.00	1.00	1.00	1066
7	1.00	1.00	1.00	748
8	1.00	1.00	1.00	358
avg / total	1.00	1.00	1.00	6115

```
[[ 919    0    0    0    0    0    0    0]
 [    0 1098    0    0    0    0    0    0]
 [    0    0  610    0    0    0    0    0]
 [    0    0    0  603    0    0    0    0]
 [    0    0    0    0  713    0    0    0]
 [    0    0    0    0    0 1066    0    0]
 [    0    0    0    0    0    0  748    0]
 [    0    0    0    0    0    0    0  358]]
```

1.0

```
tuned hyperparameters :(best parameters) {'C': 1, 'penalty': 'l1'}
accuracy : 1.0
```

5.1.5 Convert categorical variable, offer ids designated as labels 1 - 10) into dummy/indicator variables. This would enable us understand how the offered performed across transactions and demography.

```
In [97]: #Clean the income columns by filling Nan values with mean of the incomes
merged_df['income'] = merged_df['income'].fillna(merged_df['income'].mean())
X = merged_df.iloc[:,3:] #get all rows that represent quantitative data
y = merged_df.offer_successful #get all rows for offer_successful status
```

```
X = pd.concat([X, merged_df['offer_id']],axis=1)
```

```
In [98]: dummies = pd.get_dummies(X['offer_id'])
adjusted = X.drop(['offer_id', 'gender', 'became_member_on'],axis=1)
merged_new = pd.concat([adjusted, dummies],axis=1)
```

```
In [99]: merged_new.head()
```

```
Out[99]:
```

	difficulty	duration	reward	email	mobile	social	web	offer_bogo	\
0	10	7	2	1	1	0	1	0	
1	0	4	0	1	1	0	1	0	
2	0	3	0	1	1	1	0	0	
3	5	5	5	1	1	1	1	1	
4	10	10	2	1	1	1	1	0	


	offer_discount	offer_informational	...	1	10	2	3	4	5	6	7	8	9
0	1		0 ...	0	0	0	0	1	0	0	0	0	0
1	0		1 ...	0	0	0	0	0	0	0	0	0	1
2	0		1 ...	0	0	0	1	0	0	0	0	0	0
3	0		0 ...	1	0	0	0	0	0	0	0	0	0
4	1		0 ...	0	1	0	0	0	0	0	0	0	0

```
[5 rows x 26 columns]
```

```
In [103]: plt.figure(figsize=(10,10))
           # calculate the correlation matrix
           corr = merged_new.corr()

           # plot the heatmap
           #https://stackoverflow.com/questions/39409866/correlation-heatmap
           sns.heatmap(corr,
                        xticklabels=merged_new.columns,
                        yticklabels=merged_new.columns,
                        fmt='.1%', cmap='Blues')
```

```
Out[103]: <matplotlib.axes._subplots.AxesSubplot at 0x7efcd2640dd8>
```

```
In [104]: #From the above diagram, there is a correlation on the email attribute.
#This is because all customers that participated in the offers had their email address
#so let's drop the email field
X = merged_new.drop(['email'], axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_state =

scaler = preprocessing.MinMaxScaler()
```

```
In [105]: X_train.head()
```

```
Out[105]:
```

	difficulty	duration	reward	mobile	social	web	offer_bogo	\
15937	10	7	2	1	0	1	0	
15639	10	7	10	1	1	0	1	

25222	10	7	2	1	0	1	0
13047	10	7	2	1	0	1	0
2561	10	5	10	1	1	1	1

	offer_discount	offer_informational	age ...	1	10	2	3	4	5	6	7	\
15937	1	0	84 ...	0	0	0	0	1	0	0	0	
15639	0	0	118 ...	0	0	0	0	0	1	0	0	
25222	1	0	67 ...	0	0	0	0	1	0	0	0	
13047	1	0	56 ...	0	0	0	0	1	0	0	0	
2561	0	0	90 ...	0	0	0	0	0	0	0	1	

	8	9
15937	0	0
15639	0	0
25222	0	0
13047	0	0
2561	0	0

[5 rows x 25 columns]

```
In [91]: X_train = scaler.fit_transform(X_train)

X_test = scaler.fit_transform(X_test)
logistic_model = LogisticRegression(solver='liblinear',random_state=42)

In [92]: logistic_model.fit(X_train,y_train)
logistic_y_pred = logistic_model.predict(X_test)
logistic_accuracy = accuracy_score(y_test,logistic_y_pred)

print("Logistic Regression Accuracy: %.2f" % accuracy_score(y_test, logistic_y_pred))


Logistic Regression Accuracy: 0.76

In [93]: #Calculate the confusion matrix
confusion = confusion_matrix(y_test,logistic_y_pred)

In [116]: print(confusion/np.sum(confusion))

[[ 0.49178385  0.12055617]
 [ 0.1224522  0.26520777]]

In [114]: sns.heatmap(confusion/np.sum(confusion),annot=True, fmt='.1%', cmap='Blues')
plt.title('Plot of Confusion Matrix based on Logistic Regression')
plt.ylabel("Actual value")
plt.xlabel("Predicted value");
```



```
In [96]: #calculate the F1 Score
         logistic_f1 = f1_score(y_test,logistic_y_pred)
         print("F1 Score for the Logistic Regression modle is: %.2f" % logistic_f1)
```

F1 Score for the Logistic Regression modle is: 0.69

Reapply the GridSearchCV Algorithm based on the Logistic Regression trained data

```
In [106]: grid_cv2 = GridSearchCV(LogisticRegression(C=1.0, class_weight=None, dual=False, fit_i
         intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
         penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
         verbose=3, warm_start=False), params, verbose=3, n_jobs=-1,cv=3)
```

```
In [109]: grid_cv2.fit(X_train, y_train)
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

[CV] C=1, penalty=l1 ...

[LibLinear][CV] ... C=1, penalty=l1, score=0.7558807844996148, total= 0.4s

[CV] C=1, penalty=l1 ...

[LibLinear]

[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.4s remaining: 0.0s

```

[CV] ... C=1, penalty=l1, score=0.7519108846358951, total= 0.3s
[CV] C=1, penalty=l1 ...
[LibLinear]

[Parallel(n_jobs=-1)]: Done 2 out of 2 | elapsed: 0.6s remaining: 0.0s

[CV] ... C=1, penalty=l1, score=0.7453187959232046, total= 0.5s
[CV] C=1, penalty=l2 ...
[LibLinear][CV] ... C=1, penalty=l2, score=0.592937133376785, total= 0.2s
[CV] C=1, penalty=l2 ...
[LibLinear][CV] ... C=1, penalty=l2, score=0.5902707827220478, total= 0.3s
[CV] C=1, penalty=l2 ...
[LibLinear][CV] ... C=1, penalty=l2, score=0.5916093861104527, total= 0.3s
[CV] C=10, penalty=l1 ...
[LibLinear][CV] ... C=10, penalty=l1, score=0.7559992889731587, total= 0.4s
[CV] C=10, penalty=l1 ...
[LibLinear][CV] ... C=10, penalty=l1, score=0.7512591100314037, total= 0.3s
[CV] C=10, penalty=l1 ...
[LibLinear][CV] ... C=10, penalty=l1, score=0.7455558189144347, total= 0.2s
[CV] C=10, penalty=l2 ...
[LibLinear][CV] ... C=10, penalty=l2, score=0.592937133376785, total= 0.2s
[CV] C=10, penalty=l2 ...
[LibLinear][CV] ... C=10, penalty=l2, score=0.5902707827220478, total= 0.3s
[CV] C=10, penalty=l2 ...
[LibLinear][CV] ... C=10, penalty=l2, score=0.5916093861104527, total= 0.3s
[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=0.7545179830538603, total= 0.2s
[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=0.7520886413462108, total= 0.2s
[CV] C=100, penalty=l1 ...
[LibLinear][CV] ... C=100, penalty=l1, score=0.7453780516710121, total= 0.4s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.592937133376785, total= 0.2s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.5902707827220478, total= 0.3s
[CV] C=100, penalty=l2 ...
[LibLinear][CV] ... C=100, penalty=l2, score=0.5916093861104527, total= 0.3s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=0.7559400367363868, total= 0.4s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=0.7514961189784914, total= 0.4s
[CV] C=1000, penalty=l1 ...
[LibLinear][CV] ... C=1000, penalty=l1, score=0.7452002844275895, total= 0.5s
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.592937133376785, total= 0.3s
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.5902707827220478, total= 0.3s

```

```
[CV] C=1000, penalty=l2 ...
[LibLinear][CV] ... C=1000, penalty=l2, score=0.5916093861104527, total= 0.3s
[LibLinear]
```

```
[Parallel(n_jobs=-1)]: Done 24 out of 24 | elapsed: 7.3s finished
```

```
Out[109]: GridSearchCV(cv=3, error_score='raise',
                      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                      intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                      penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                      verbose=3, warm_start=False),
                      fit_params=None, iid=True, n_jobs=-1,
                      param_grid={'penalty': ['l1', 'l2'], 'C': [1, 10, 100, 1000]},
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring=None, verbose=3)
```

```
In [110]: grid_cv2.best_params_
```

```
Out[110]: {'C': 1, 'penalty': 'l1'}
```

```
In [111]: logistic_y_pred2 = grid_cv2.predict(X_test)
          logistic_accuracy2 = accuracy_score(y_test,logistic_y_pred2)
          logistic_f12 = f1_score(y_test,logistic_y_pred2)
```

```
In [112]: #Generate the Confusion matrix
          conf_matrix2 = confusion_matrix(y_test, logistic_y_pred2)
```

```
print("See the classification report below:")
print(classification_report(y_test, logistic_y_pred2))
print("See the confusion matrix below:")
print(conf_matrix2)
print("The accuracy from Logistic Regression using Grid Search CV is: ", logistic_accuracy2)

print("Logistic Regression : (best parameters) ", grid_cv2.best_params_)
print("Logistic Regression accuracy : ", grid_cv2.best_score_)
```

```
See the classification report below:
```

	precision	recall	f1-score	support
0.0	0.80	0.80	0.80	7751
1.0	0.69	0.68	0.69	4907
avg / total	0.76	0.76	0.76	12658

```
See the confusion matrix below:
```

```
[[6224 1527]
 [1555 3352]]
```

The accuracy from Logistic Regression using Grid Search CV is: 0.756517617317
Logistic Regression :(best parameters) {'C': 1, 'penalty': 'l1'}
Logistic Regression accuracy : 0.751036934624

```
In [295]: #Confusion Matrix expressed as a percentage
          percent_conf = conf_matrix2/np.sum(conf_matrix2)
```

```
In [298]: print("See the confusion matrix below: (%)")
          print(percent_conf)
```


```
See the confusion matrix below: (%)
[[ 0.49138884  0.12095118]
 [ 0.1224522  0.26520777]]
```

```
In [299]: print("Logistic Regression FI Score :",logistic_f12)
```

Logistic Regression FI Score : 0.685451761103

```
In [113]: sns.heatmap(conf_matrix2/np.sum(conf_matrix2),annot=True, fmt='.1%', cmap='Blues')
          plt.title('Plot of Confusion Matrix based on GridSearchCV with Logistic Regression')
          plt.ylabel("Actual value")
          plt.xlabel("Predicted value");
```

Plot of Confusion Matrix based on GridSearchCV with Logistic Regression



```
In [ ]:
```